

PARALLEL IMPLEMENTATION OF HASHING ALGORITHMS

PROJECT REPORT FOR CS359

Submitted by

Kunal Gupta (150001015)

Keshav Goyal (150001014)

Computer Science and Engineering

3rd Year

Under the Guidance of

Dr. Surya Prakash



Department of Computer Science and Engineering

Indian Institute of Technology Indore

Autumn 2017

CONTENTS

INTRODUCTION	3
PROCEDURE	3
<i>Selection of Hashing Algorithm</i>	<i>3</i>
<i>Selection of Bucketing and Hashing Functions.....</i>	<i>4</i>
<i>Implementation</i>	<i>5</i>
<i>Analysis.....</i>	<i>5</i>
<i>Detecting and Removing Problems.....</i>	<i>5</i>
ALGORITHMS.....	5
<i>Sequential Cuckoo Hash.....</i>	<i>5</i>
<i>Parallel Cuckoo Hash</i>	<i>6</i>
OBSERVATIONS AND ANALYSIS.....	7
<i>Effect of Input Size on Execution Time and Speedup.....</i>	<i>7</i>
<i>Effect of table size on Number of Collisions</i>	<i>9</i>
<i>Optimal Table Size</i>	<i>11</i>
RESULTS	11
CONCLUSIONS	11
DETECTING AND REMOVING PROBLEMS.....	12
<i>Cycles.....</i>	<i>12</i>
<i>Theoretical Table Size</i>	<i>12</i>

INTRODUCTION

With increase in size and complexity of data, the storage and access of data is an important task. One way in which data is structured is in the forms of pairs, like Product-Price in an e-commerce business, Product-Stock in inventory systems, etc. Here the insertion, access and updating of the data can be very costly for large dataset size. This pair like association is generally implemented with hash tables/maps. Hashing is one of the methods used to map such key-value-like data.

In this project, we concentrate on the problem of implementing a parallel-friendly data structure that allows efficient random access of millions of elements and can be both constructed and accessed at interactive rates. Such a data structure has numerous applications in computer graphics, centered on applications that need to store a sparse set of items in a dense representation.

Cuckoo hashing is one of the schemes for resolving hash collisions of values of hash functions in a table, with worst-case constant lookup time. It is a so-called “multiple-choice” perfect hashing algorithm that achieves high occupancy at the cost of greater construction time.

PROCEDURE

SELECTION OF HASHING ALGORITHM

There are several algorithms that are used to hash values into a Hash Table. They may be divided on the basis of their way of handling collisions, some of which are:

- Linear Probing
 - Placing the new key into the closest following empty cell.
 - Results in primary clustering.
 - As the cluster grows larger, the search for those items hashing within the cluster becomes less efficient.
- Quadratic Probing
 - Given a collision with Hash Value H , the probing sequence now generated is:

$$H + 1^2, H + 2^2, H + 3^2, H + 4^2, \dots, H + k^2$$
 - Provides good memory caching because it preserves some locality of reference.
 - There is no guarantee of finding an empty cell once the table gets more than half-full, or even before the table gets half-full if the table size is not prime.
- Cuckoo Hashing
 - Resolves collisions by using more than one (k) Hash Functions (generally, $k = 2$).
 - This provides more than one possible locations in the hash table for each key.

- Lookup requires inspection of just k locations in the hash table, which takes constant time in the worst case.
- However, *cycles* may be present.

We chose cuckoo hash as our open-addressing scheme to remove collisions as calculation of the next available location can be parallelized in a much simpler and efficient way.

SELECTION OF BUCKETING AND HASHING FUNCTIONS

In our implementation, each key is assigned a bucket. Bucketing allows for collision reduction by restricting collisions to inside the bucket. Buckets may be assigned by a function such as:

$$b(k) = k \bmod |\text{buckets}|$$

$$\text{or } b(k) = [(c_0 + c_1 \cdot k) \bmod 1900813] \bmod |\text{buckets}|$$

where c_0 and c_1 are random integers and 1900813 is a prime number.

We realized that for a bucket size s , such that $s \times |\text{buckets}| = \text{Number of keys}$, with a uniform distribution of keys in the range of 0 to N , the first hash function will suffice along with the buckets filling up uniformly. We choose $N/409$ buckets with bucket size as 409.

Now we use cuckoo hashing to hash values inside each bucket. Three hash functions will suffice, since using just two functions permits a load factor above 80% (theoretically). Algebraic hash functions of the following form are chosen:

$$g_i(k) = [(c_{i,0} + c_{i,1} \cdot k) \bmod 1900813] \bmod |T_i|$$

All constants $c_{i,j}$ are random numbers.

By noticing the above functions, it can be seen that instead of having a single hash table, three hash tables, one for each hash function, are created to minimize collisions and allow for easier parallelization.

Initially we choose Table Size $T_i = 192$, thus overall load factor of $\frac{409}{192 \times 3} = 71\%$.

IMPLEMENTATION

Initially we modified and appropriately designed the algorithm for sequential Cuckoo Hash and then the Parallel Cuckoo Hashing and then implemented the same in OpenMP.

ANALYSIS

All analysis is done on an Intel i7-4710MQ Processor (4th Generation i7 with 4 cores/8 Threads).

- Speed-Up calculation for various input sizes, keeping all other parameters fixed.
- Number of collisions due to cycles for various table size, keeping input size fixed.
- Optimum table size to accommodate all keys by rebuilding the hash tables less than 100 times.

DETECTING AND REMOVING PROBLEMS

As discussed further in the report.

ALGORITHMS

SEQUENTIAL CUCKOO HASH

The set of keys = **keys**, hashing functions = $g_i(k)$ and Hash Tables = T_i .

```

0. flag = False
1. for k  $\in$  keys:
2.    $loc_1 = g_1(k)$ 
3.   if ( $k \notin T_1$  and  $k \notin T_2$  and  $k \notin T_3$ ):
4.      $T_1[loc_1] = k$ 
5.   for k  $\in$  keys:
6.      $loc_2 = g_2(k)$ 
7.     if ( $k \notin T_1$  and  $k \notin T_2$  and  $k \notin T_3$ ):
8.        $T_2[loc_2] = k$ 
9.   for k  $\in$  keys:
10.     $loc_3 = g_3(k)$ 
11.    if ( $k \notin T_1$  and  $k \notin T_2$  and  $k \notin T_3$ ):
12.       $T_3[loc_3] = k$ 
13.   for k  $\in$  keys:
14.     if ( $k \notin T_1$  and  $k \notin T_2$  and  $k \notin T_3$ ):
15.       flag = True
16. if (flag):
17.   goto Step 0
  
```

PARALLEL CUCKOO HASH

```

0. flag = False
1. parfor k ∈ keys:
2.   loc1 = g1(k)
3.   if (k ∉ T1 and k ∉ T2 and k ∉ T3):
4.     lock(T1[loc1])
5.     T1[loc1] = k
6.     unlock(T1[loc1])
7. parfor k ∈ keys:
8.   loc2 = g2(k)
9.   if (k ∉ T1 and k ∉ T2 and k ∉ T3):
10.    lock(T2[loc2])
11.    T2[loc2] = k
12.    unlock(T2[loc2])
13. parfor k ∈ keys:
14.   loc3 = g3(k)
15.   if (k ∉ T1 and k ∉ T2 and k ∉ T3):
16.    lock(T3[loc3])
17.    T3[loc3] = k
18.    unlock(T3[loc3])
19. parfor k ∈ keys:
20.   if (k ∉ T1 and k ∉ T2 and k ∉ T3):
21.    flag = True
22. if(flag):
23.   goto Step 0

```


Here we follow the policy that the key, which collides to same value with same hash function and written earlier, is kicked out. Here, we have to lock the location before accessing; else, there will be corruption of data for obvious reasons.

OBSERVATIONS AND ANALYSIS

EFFECT OF INPUT SIZE ON EXECUTION TIME AND SPEEDUP


SMALL INPUT

INPUT SIZE (IN 10 THOUSANDS)	SEQUENTIAL EXECUTION TIME (IN SECONDS)	PARALLEL EXECUTION TIME (IN SECONDS)	SPEED-UP
1	0.014536	0.005092	2.854674
2	0.029888	0.029288	1.020486
3	0.043821	0.028131	1.557748
4	0.072915	0.031171	2.339193
5	0.078906	0.035846	2.20125
6	0.093383	0.049764	1.876517
7	0.108593	0.050698	2.141958
8	0.12498	0.048964	2.552488
9	0.145717	0.064958	2.243249
10	0.158111	0.071268	2.218541
11	0.181713	0.064454	2.819266
12	0.193379	0.080252	2.409647
13	0.207849	0.087912	2.364285
14	0.230006	0.0963	2.388432
15	0.245186	0.098222	2.496243
16	0.310394	0.117737	2.636334
17	0.280346	0.165282	1.696168
18	0.299863	0.113387	2.644598
19	0.306185	0.132157	2.316828
20	0.326376	0.112476	2.901739



LARGE INPUT

INPUT SIZE (IN 10 THOUSANDS)	SEQUENTIAL EXECUTION TIME (IN SECONDS)	PARALLEL EXECUTION TIME (IN SECONDS)	SPEED-UP
10	0.157277	0.05715	2.752003
15	0.242187	0.07282	3.325831
20	0.331504	0.097005	3.417391
25	0.417652	0.129555	3.223743
30	0.504879	0.145762	3.463722
35	0.620299	0.179983	3.446431
40	0.712621	0.193404	3.684624
45	0.808007	0.22057	3.663268
50	0.949989	0.24317	3.906687
55	1.072394	0.261115	4.10698
60	1.280823	0.315909	4.054405
65	1.387326	0.340306	4.076702
70	1.468277	0.349124	4.205603
75	1.657462	0.413346	4.009866
80	1.732724	0.444641	3.896906
85	2.403379	0.495043	4.854889
90	2.244876	0.510905	4.393921
95	2.506213	0.537177	4.665526
100	3.177639	0.57041	5.570798




Large input size is thus better for performance analysis. We observe that speedup increases with the size of the input; and for a 10-fold increase in input size, sequential time increases by 20 times whereas, parallel time by only a 10-fold.

The execution time is linearly dependent on the input size as evident from the graph.


EFFECT OF TABLE SIZE ON NUMBER OF COLLISIONS

Here we ran the algorithm for 100 iterations for each value of Hash Table Size and calculated average number of collisions for N=100,000. We repeated this 4 times:

TABLE SIZE	AVERAGE NUMBER OF COLLISIONS	1	2	3	4
190	5805	5971	5886	5444	5918
200	4769	4386	5234	4609	4848
210	4107	4228	4078	4057	4063
220	3089	2580	3322	2779	3675
230	2842	2397	2700	2982	3290
240	2331	3128	2024	2015	2157
250	2008	1799	1976	2120	2138
260	1830	1747	2470	1710	1394
270	1587	2135	1487	1346	1380
280	1583	1387	1718	1612	1616
290	1305	1110	1941	1251	917
300	1132	773	1167	1064	1525
310	1051	760	1051	1343	1051
320	666	688	667	703	606
330	651	871	490	610	631
340	683	1123	555	413	639
350	659	557	548	803	727
360	592	502	800	681	386
370	473	476	366	543	505
380	389	427	324	553	252
390	389	303	465	426	362
400	280	276	175	396	273
410	442	533	335	299	600
420	342	249	355	415	350
430	329	238	447	306	323
440	276	164	226	315	400
450	227	317	137	138	317
460	223	291	196	251	152
470	286	183	181	637	144
480	201	313	130	211	150
490	126	77	168	99	159
500	137	174	192	88	95
510	140	82	148	278	51
520	201	162	100	177	364
530	151	110	205	128	160
540	115	105	86	73	195
550	106	141	54	79	150
560	71	90	75	59	58
570	141	130	196	147	91
580	61	37	76	66	65
590	64	95	41	82	39
600	114	47	44	152	214



Since the trends in individual observations have many spikes, an average of the four observations will provide with better trends of the number of collisions with table size:



We observe that merely doubling the table size reduces the number of collisions to 0.16% of its initial value.

OPTIMAL TABLE SIZE

From the last observations, we also noted the least table size that was required for the hash to be built in no more than 100 iterations of rebuilding (i.e. re-running cuckoo-hash) and selecting a new hash functions:

OBSERVATION NUMBER	1	2	3	4
TABLE SIZE	320	330	270	310

We take average of the least table size, $T_i \approx 307 \pm 22$ (*with load factor = 44%*) such that we always have a good chance of building the hash table in no more than 100 iterations of rebuilding the Hash Tables and regenerating new Hash Functions. This can further be checked for larger number of iterations of rebuilds to have a decreased table size and a load factor more than 44%.

RESULTS

- The dependence of Execution Times on Input size seems to be linear.
- Speedup increased with increase in input size.
- The increase in the size of Hash Table led to a drastic decrease in the number of collisions.
- In order to get a good chance of accommodating the Hash Tables in less than 100 iterations of re-building the Hash Table, a good table size is $T_i = 307 \pm 22$ (*with load factor = 44%*)

CONCLUSIONS

Cuckoo Hash is a highly parallelizable algorithm reaching load factors of upto 90% (theoretical). For less than 100 rebuilds of the Hash Tables, a minimum of load factor = 44% for input of 100,000 keys.

The above results depict that its parallel algorithm outperforms the sequential algorithm by a considerable margin. The difference in performance is even larger when input size increases.

DETECTING AND REMOVING PROBLEMS

CYCLES

A cycle is kicking-out of elements in a cyclic fashion across hash tables. These cycles will lead to infinite loops. Detection of the cycle can be done by counting the number of times we had to perform kicks for all collision values and if it is more than a threshold value, we can assume that a cycle has been created. We can then regenerate the hash functions and start hashing from the scratch.

We tested the threshold to be equal to 25 and observed that increasing it did not lead to decreased number of collisions, thus it being a satisfactory value for determining a cycle.

THEORETICAL TABLE SIZE

The theoretically optimum table size with a load factor of 71% did not work for small number of iterations of the algorithm, thus we looked at different table sizes and their efficiency of hashing in under 100 iterations, although there might be a dependence on the pseudo random generator of C++ (which we ignored).